

Web Databases (SQL and NoSQL)

Lecture 11 (CPEN 400A)
Karthik Pattabiraman

Based on CS498RK at UIUC (used with permission),
and the MongoDB tutorial (docs.mongodb.)

Outline

- Relational Databases (SQL-based)
- ACID semantics
- Non-traditional Databases (NoSQL)
- MongoDB Primer

What's a Database ?

- In its simplest form, it's a collection of data
 - Allows applications to modify/access data through standard interfaces
 - Separate data storage from logical organization
- Many types of databases
 - Hierarchical
 - Object oriented
 - **Relational**
 - Document-based

↳ NoSQL

→ Most popular (SQL - current)
- Modelled mathematically
- Many optimizations modelled
- used by many companies, oracle.

Relational Database

- Stores the data in the form of tables (Relations) to map one kind of data to another
- Why tables ?
 - Separate data storage from logical view of data
 - Easy to express relationships between data
 - Aggregate data from multiple tables on demand (table joins)
 - Allow declarative queries to be executed

Example of a Table

- Much like a spreadsheet, except the columns are of fixed type and rows are identified by a unique key (known as primary key)

id	given_name	middle_name	family_name	date_of_birth	grade_point_average	start_date
1	Giles	Prentiss	Boschwick	3/31/1989	3.92	9/12/2006
2	Milletta	Zorgos	Stim	2/2/1989	3.94	9/12/2006
3	Jules	Bloss	Miller	11/20/1988	2.76	9/12/2006
4	Greva	Sortingo	James	7/14/1989	3.24	9/12/2006
...

↳ unique Key

Source:

<http://archive.oreilly.com/pub/a/ruby/excerpts/ruby-learning-rails/intro-ruby-relational-db.html>

- Rigid col^c types , exact format of types

Database schema

- A logical representation of the tables' structure listing each column name and type

Column Name	Type
id	Integer
given_name	String
middle_name	String
family_name	String
date_of_birth	Date
grade_point_average	Floating Point
start_date	Date

- No SQL do not necessarily have schema.

Multiple Unconnected Tables

id	given_name	middle_name	family_name	date_of_birth	grade_point_average	start_date
1	Giles	Prentiss	Boschwick	3/31/1989	3.92	9/12/2006
2	Milletta	Zorgos	Stim	2/2/1989	3.94	9/12/2006
3	Jules	Bloss	Miller	11/20/1988	2.76	9/12/2006
4	Greva	Sortingo	James	7/14/1989	3.24	9/12/2006
...

id	username	password_hash	role
763	Demetrius	ASVUQP8AZV8	administrator
845	Sharon	8WEROCPA387	class_admin
973	Wilmer	S3D03VP3A8AS	class_admin
1021	Nicolai	SDF83NC9A2F2J	data_analyst

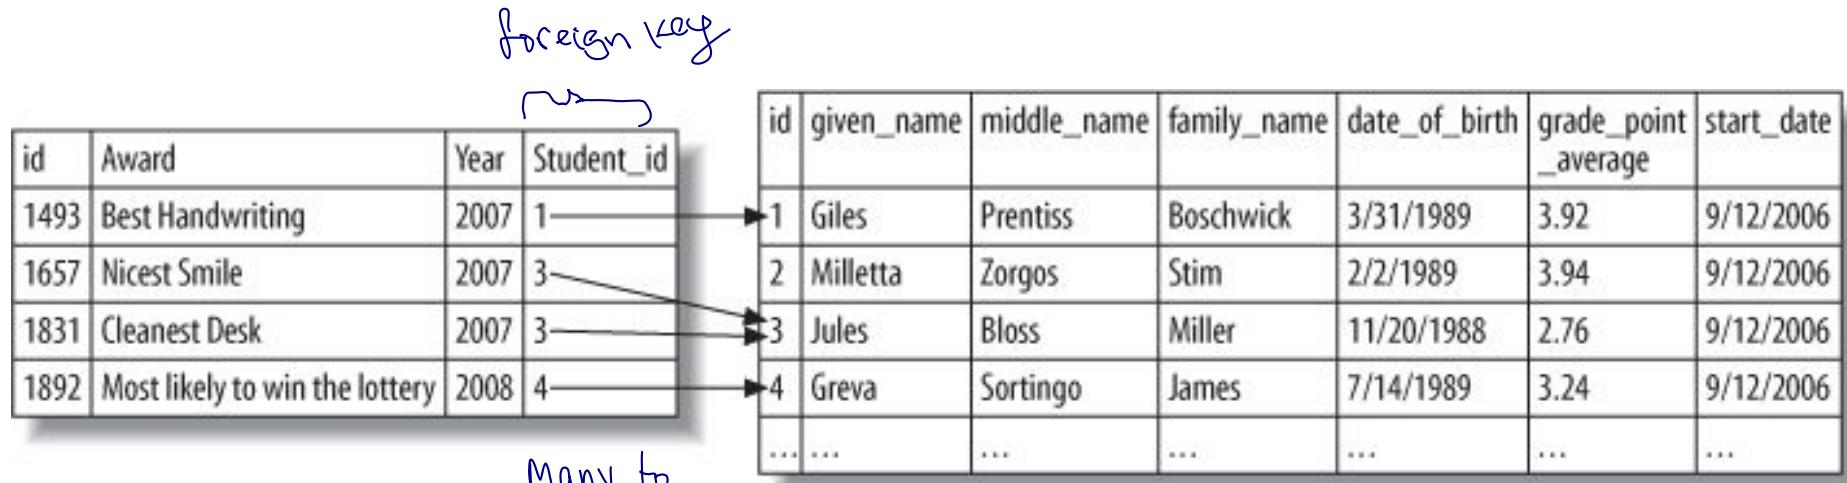
Source:

<http://archive.oreilly.com/pub/a/ruby/excerpts/ruby-learning-rails/intro-ruby-relational-db.html>

Connected Tables

- The problem with having multiple unconnected tables is that it's difficult to tell if the same record is present in both tables
 - **Solution 1 (Ugly):** Duplicate the relevant data in each table. Complicates data management, updates and need to anticipate queries in advance
 - **Solution 2 (Preferred):** Keep a pointer (foreign key) to the other table so that you can access the data by following the pointer. No need to anticipate queries in advance, easy to modify

Connected Tables



Source:

<http://archive.oreilly.com/pub/a/ruby/excerpts/ruby-learning-rails/intro-ruby-relational-db.html>

Each table has what is known as **primary key** to uniquely identify records in it.

Tables keep **foreign keys** to link to records in other tables. A foreign key is the primary key of the table being linked to.

Table Joins

- Can be used to combine information from multiple tables together (e.g., through SQL)
 - Produces a single table containing the information in both tables, without **duplication**
 - Joins can involve more than one table
- For example, we can produce a single join table having the award name and the student details from the previous slide

Example of a Join in SQL

- SELECT * from Employees, Departments
where employee.deptID=department.deptID

Employee table		Department table	
LastName	DepartmentID	DepartmentID	DepartmentName
Rafferty	31	31	Sales
Jones	33	33	Engineering
Heisenberg	33	34	Clerical
Robinson	34	35	Marketing
Smith	34		
Williams	NULL		



list everything
together →

Employee.LastName	Employee.DepartmentID	Department.DepartmentName	Department.DepartmentID
Robinson	34	Clerical	34
Jones	33	Engineering	33
Smith	34	Clerical	34
Heisenberg	33	Engineering	33
Rafferty	31	Sales	31

The problem with Joins

→ e.g. 10 ways join, w/ distributed sys
→ need to be avail, all time to join

- Joins are expensive as they need to straddle multiple tables (potentially stored elsewhere)
- Combination of fields from different tables can result in losing cache locality
- Join performance is poor for large tables, though databases are very good at optimizing
- Requires all the tables to be available during join - otherwise join will fail (more later)

↗ recommend min hw configuration

Outline

- Relational Databases (SQL-based)
- **ACID semantics**
- Non-traditional Databases (NoSQL)
- MongoDB Primer

SQL Databases have ACID Semantics

↳ Satisfy all properties

ATOMICITY *all or nothing*

CONSISTENCY *written data follows rules and constraints*

ISOLATION *uncommitted transactions are isolated from each other*

DURABILITY *committed transactions are permanent*

Atomicity and Transactions

- Transaction is a sequence of operations which are executed all at once or not at all
(Atomicity) → cannot have partial completion
 - If failures occur, roll-back to the beginning
 - **Example:** Transfer \$1000 from Accts. A to B
 - Step 1: Locate Account A and check balance
 - Step 2: Subtract 1000 dollars from Acct A
 - Step 3: Credit 1000 dollars to Acct B
- All parts needed to complete*
- roll transaction back if fail, commit if complete
- } wrap in 1 transaction.*

Consistency

- Can check one or more constraints on the resulting data, and abort if not satisfied

```
CREATE TABLE acidtest (  
    A INTEGER, B INTEGER,  
CHECK (A + B = 100));
```

- invariants should be maintained
- checks written by programmers, rollback if fail.

Isolation

- Transactions are isolated from one another

*transactions
should not cross - each should
have own "view"*

~~X~~ **T1 subtracts 10 from A**

T2 subtracts 10 from B

T2 adds 10 to A

T1 adds 10 to B

- Support large N^{∞} of transaction

Durability

- Transactions are permanent when committed

T1 subtracts 10 from A

T1 adds 10 to B

T2 subtracts 10 from B

T2 adds 10 to A

- from commit, cannot change record-
- if non permanent, cannot make progress

ACID: Pros and Cons

- **Pros**

- Simplifies reasoning about actions of the system
- Guarantees correctness in presence of failures

↳ Cannot guarantee for NoSQL database.

- **Cons**

- Guarantees come with huge performance cost
- Cannot guarantee availability when network fails
 - This is due to something called the CAP theorem

• old relDB, will stick to each other / tightly connected
• lockstep w/ each other
– O.K. if data is important (e.g. bank) → Web: handling failure more important.

Class Activity

- Consider the following transactions T1 and T2 which execute on a bank account database. Which of the four ACID rules, if any, (Atomicity, Consistency, Isolation, Durability) are violated ?
- Assume initial balance is \$100. T1 attempts to deposit \$900 to the account. At the same time, T2 checks if the account balance ≥ 500 and returns true. However, T1 aborts and the account balance becomes \$100 again.

Outline

- Relational Databases (SQL-based)
- ACID semantics
- **Non-traditional Databases (NoSQL)**
- MongoDB Primer

↗ "Not SQL" original

↓ now

Not only SQL

Can support Obj model2.

NoSQL Databases

- Do not natively support Table joins
 - Are much more scalable and failure tolerant
 - Must do joins explicitly using program code
 - Do not typically support ACID semantics
 - So data may be inconsistent or out of sync
(provide what is known as eventual consistency)
 - When failures occur, data may be lost or incorrect
 - Failure tolerant, but
- if you wait long enough, it might be consistent*

CAP Theorem [Brewer'99]

- You can achieve only two of the following three properties in any database system

CONSISTENCY "...requiring requests of the distributed shared memory to act as if they were executing on a single node, responding one at a time" - Cannot have two req. return inconsistent val.

AVAILABILITY "... every request received by a non-failing node in the system must result in a response" - reasonable response time. Pick 'of 2'

PARTITION TOLERANCE "... the network will be allowed to lose arbitrarily many messages sent from one node to another"

L) Network partition \Rightarrow nodes cannot comm. \rightarrow 1. either respond
2. wait. \Rightarrow Reality Nat option.

CAP theorem continued..

- During a network partition, a system must choose either consistency or availability for it to work through the partition
 - Traditional SQL-based databases choose consistency and may hence not be available
 - NoSQL databases choose availability and hence may not be consistent
 - In web applications, availability often trumps consistency

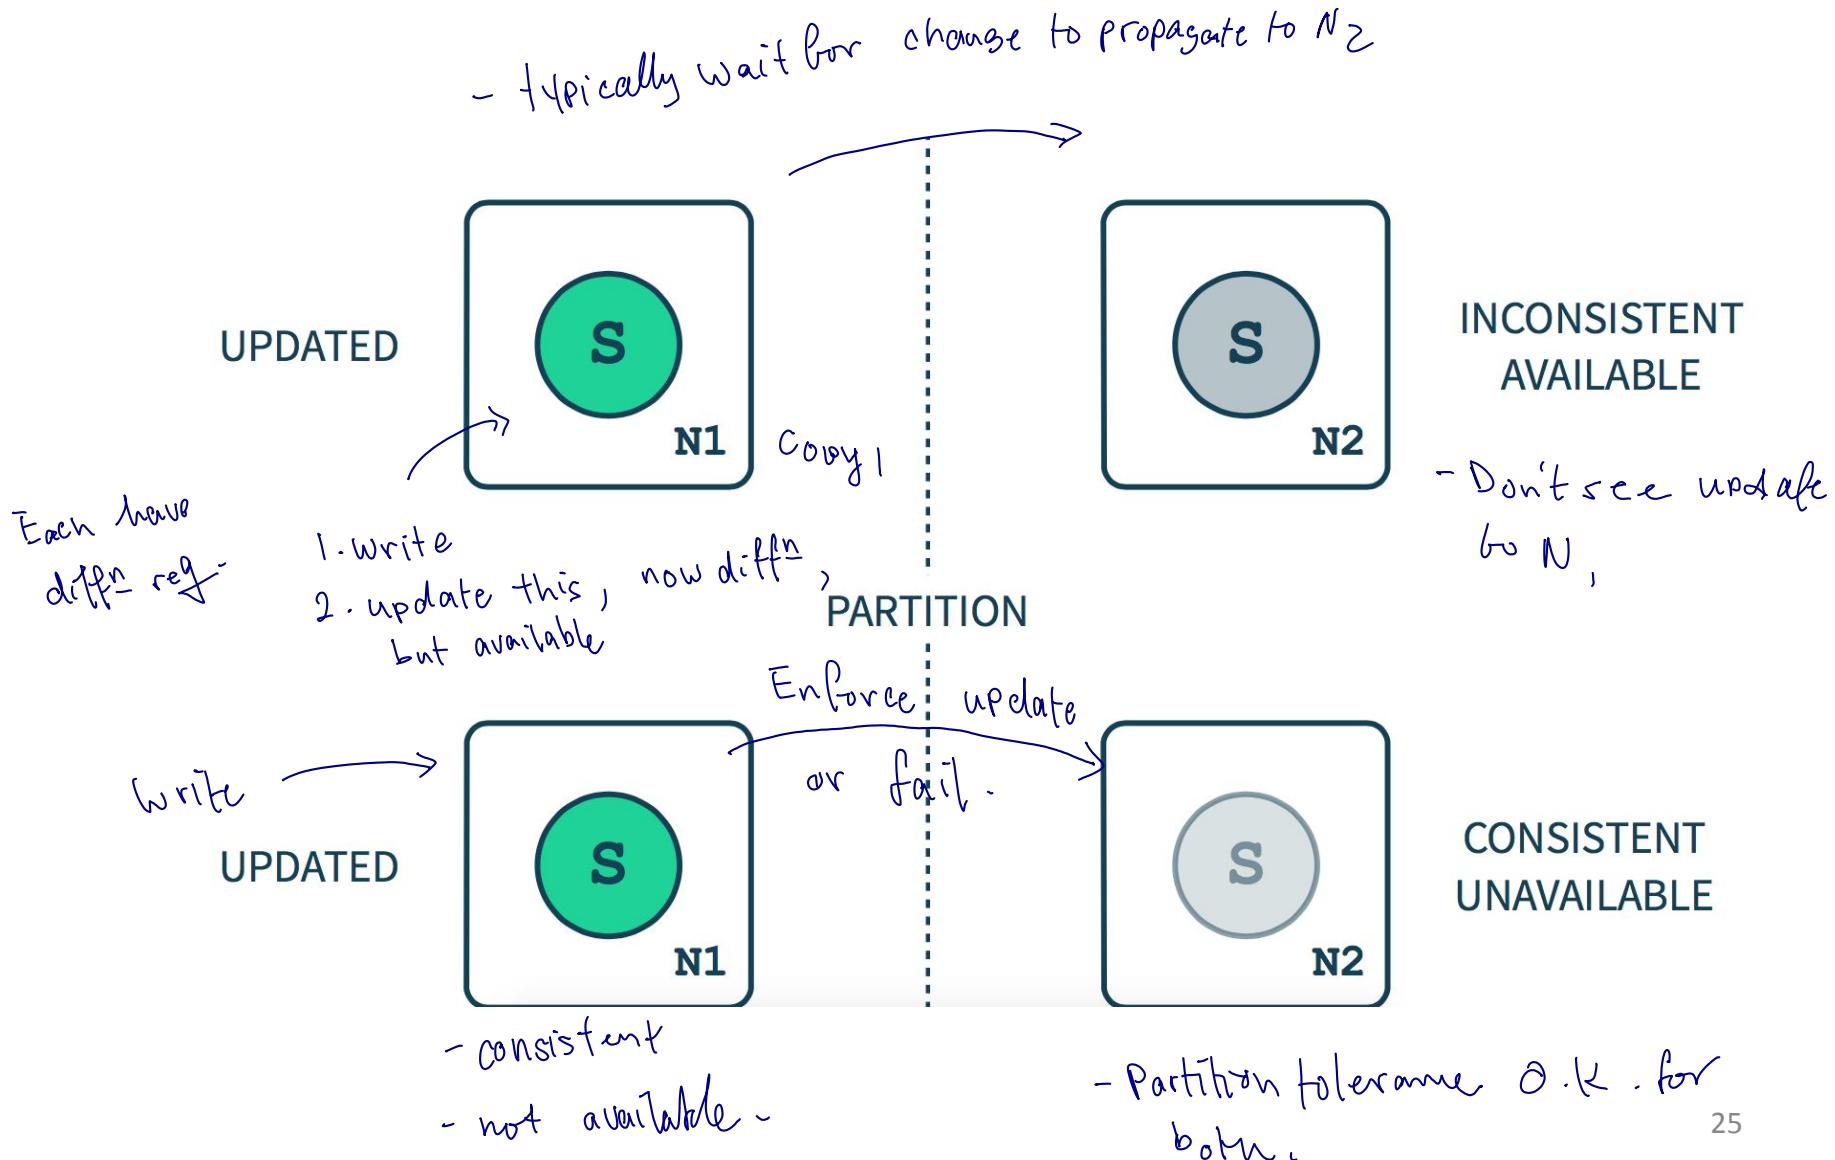
↳ e.g. facebook feed.

- doesn't matter if it is not the same

bank
- consistency
- do not care if ATM out of serv.²⁴

- only applies to Write + read. W/ read, never inconsistent.

Example of Network Partitioning

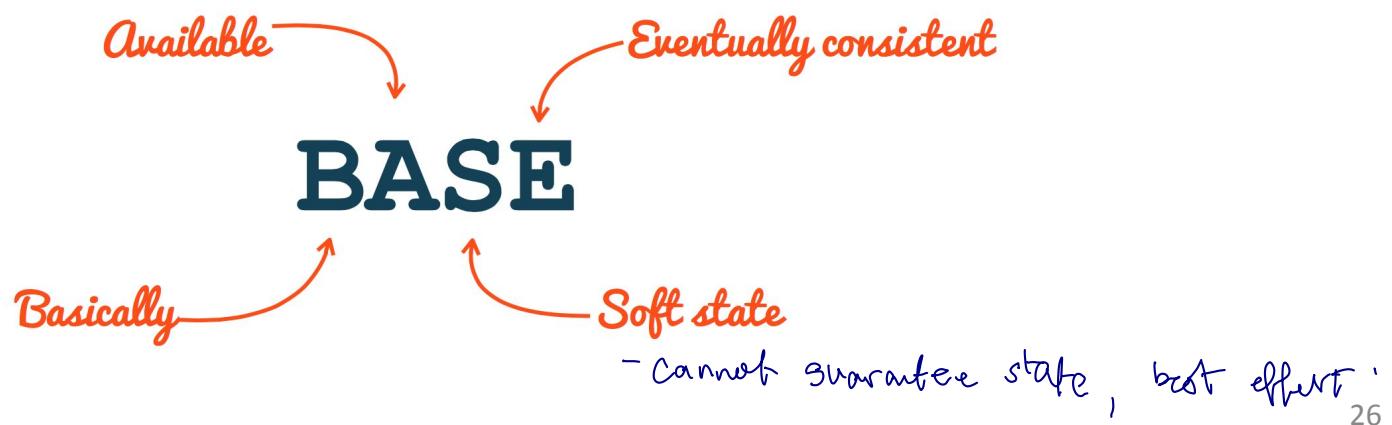


Eventual Consistency

- NoSQL databases provide a guarantee that they will eventually be consistent (e.g., when the network partition heals)
 - Eventually can be a very long time
 - Consistent does not mean correct....
 - things happen in between compromising correctness.

→ when partition heals.

→ no guarantee



SQL Vs NoSQL - 1

SQL

TYPES

one type

- relational, typically

EXAMPLES

MySQL, SQLite,
Oracle Database

NoSQL

key-value,
document, graph

MongoDB,
Cassandra, HBase,
Neo4j

SQL Vs. NoSQL - 2

Glorified hashtable
↑

SQL

DATA STORAGE MODEL

Individual records are stored as rows; columns store a specific piece of data about record

Separate data types are stored in separate tables and joined together when complex queries are executed

NoSQL

Key-value stores are similar to SQL, but have only two columns

Document DBs store all relevant data together in a single document in a hierarchically nested format (JSON, XML)

www.mongodb.com/nosql-explained

SQL Vs. NoSQL - 3

SQL

SCHEMAS

Structure and data types are fixed in advance

SCALING

Vertically: single server must be made increasingly powerful

- runs multiple close, lockstep

NoSQL

Unlike SQL rows, dissimilar data can be stored together as necessary

- skip schema, since it is only for table join.

- Horizontally: distribute data over several machines

- data center model.

SQL Vs. NoSQL - 4

SQL

**SUPPORTS
TRANSACTIONS**

Yes

CONSISTENCY

Strong consistency

You can set
consistency level

NoSQL

In certain circumstances
and at certain levels
(document-level) ↗
*transaction is state of all.
single record consistent.*

Tunable consistency
(MongoDB), Eventual
consistency (Cassandra)

Class Activity

- For each of the following scenarios, will you use a traditional database or non-SQL database. Justify your answer using CAP thrm.
 - Online photo gallery to browse photos and upload photos occasionally from multiple locations
 - Large ecommerce store in which the inventory needs to reflect any purchases made instantly in all locations
 - Shopping cart of customers in an online store in which users can login from different locations

Outline

- Relational Databases (SQL-based)
- ACID semantics
- Non-traditional Databases (NoSQL)
- **MongoDB Primer**

MongoDB

- **Document-oriented NoSQL database**
 - Documents are the equivalent of tables
 - Stored in JSON format (technically BSON, or binary JSON)
 - Must be smaller than 16 MB in size
 - other weird restrictions
- **No apriori schema needed, or rather schema can be modified dynamically**
 - Can store dissimilar objects in same document
 - Documents can be embedded in other documents
 - MongoDB supports JS. natively, pass JSON back & forth

MongoDB: Data types

JSON: null, boolean, number, string, array, and object

MongoDB: null, boolean, number, string, array, **date, regex, embedded document, object id, binary data, code**

MongoDB: Example Dataset

```
{  
  "address": {  
    "building": "1007",  
    "coord": [ -73.856077, 40.848447 ],  
    "street": "Morris Park Ave",  
    "zipcode": "10462"  
  },  
  "borough": "Bronx", ← Cat  
  "cuisine": "Bakery",  
  "grades": [ ← Rrr.  
    { "date": { "$date": 1393804800000 }, "grade": "A", "score": 2 },  
    { "date": { "$date": 1378857600000 }, "grade": "A", "score": 6 },  
    { "date": { "$date": 1358985600000 }, "grade": "A", "score": 10 },  
    { "date": { "$date": 1322006400000 }, "grade": "A", "score": 9 },  
    { "date": { "$date": 1299715200000 }, "grade": "B", "score": 14 }  
  ],  
  "name": "Morris Park Bake Shop",  
  "restaurant_id": "30075445"  
}
```

Databases and Collections

- A MongoDB database consists of multiple databases. Specify db to use by “use test”
z → Name of DB.
- A database can have multiple collections. Specify collection as db.collectionName.op
z → operator.
- A collection can have one or more documents
 - Each record is called a document

Insert into a Database

- db.collectName.insert(document in JSON)

The diagram shows a MongoDB insert statement with handwritten annotations:

```
db.restaurants.insert(  
  {  
    "address" : {  
      "street" : "2 Avenue",  
      "zipcode" : "10075",  
      "building" : "1480",  
      "coord" : [ -73.9557413, 40.7720266 ],  
    },  
    "borough" : "Manhattan",  
    "cuisine" : "Italian",  
    "grades" : [  
      {  
        "date" : ISODate("2014-10-01T00:00:00Z"),  
        "grade" : "A",  
        "score" : 11  
      },  
      {  
        "date" : ISODate("2014-01-16T00:00:00Z"),  
        "grade" : "B",  
        "score" : 17  
      }  
    ],  
    "name" : "Vella",  
    "restaurant_id" : "41704620"  
  })
```

- Collection name: A blue bracket points from the word "restaurants" to the text "collection name".
- Method: A blue bracket points from the ".insert()" method to the text "method".
- Restaurant in JSON: A blue bracket points from the entire document to the text "Restaurant in JSON".
- J.S. code.: A blue bracket points from the first part of the list to the text "- J.S. code .".
- typically use API for db application: A blue bracket points from the second part of the list to the text "- typically use API for db application".
- or query db directly: A blue bracket points from the third part of the list to the text "- or query db directly".

Finding objects

- `db.collectName.find()` – shows all documents
- `db.collectName.find(JSON object)` – shows documents satisfying the given JSON object
 - Finds all docs with the fields and values equal to the JSON object passed as an argument
 - Can also specify conditional operations such as `$lt`, `$gt`, or logical combinations (using AND, OR)

Examples of queries

- db.restaurants.find({"borough": "Manhattan"})
 - Finds all restaurants with the borough==manhattan

→ JSON obj.

- db.restaurants.find({ "grades.score": { \$gt:30} })

Find restaurants wl grade.score greater than 30.

Object_id

- Every document is given a unique ‘_id’ value – automatically assigned by the MongoDB
- Object IDs must be unique in a document, and should be of type ObjectId
- Can be used to remove or update specific objects

Update

- db.collectName.update(objects to be matched, object fields to be updated)

```
db.restaurants.update(  
  { "name" : "Juni" },   ← match  
  {  
    $set: { "cuisine": "American (New)" },  
    $currentDate: { "lastModified": true }  
  }  
)
```



← update operator

Update operator (full list of operators can be found at:
<https://docs.mongodb.org/manual/reference/operator/update/>)

Remove

- Can remove documents from a collection using the remove method

`db.collectName.remove(matching condition)`

example: `db.restaurants.remove({ "borough": "Manhattan" })`

– remove restaurants in Manhattan

Operations on each record

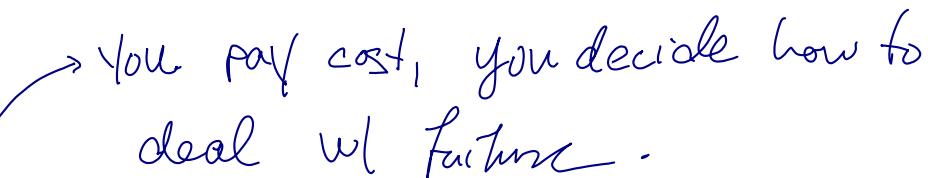
Example: Print the grades of all restaurants that have more than one grade associated with them.

- iter - e thru

```
db.restaurants.find().forEach(  
    function(Object) {  
        if (Object.grades.length > 1)  
            printjson(Object.grades);  
    }  
)
```

-JS direct exec on dB ,

Table Joins in MongoDB

- Joins are not natively supported in MongoDB and hence need to be written manually
 - Iterate over each document of the first collection
 - Lookup the corresponding document in the second collection either by key or by name
 - Write JavaScript code to merge the information in the relevant fields from the two documents
 - Return the merged information as the query result
- No free lunch - 

Example: Join Operation

- Assume that you had another collection in the database called “users” which had a list of users who had reviewed each restaurant. Assume this collection is indexed by restaurant name.
- We wish to write a query to list all the restaurants that have at least one review, and the list of users who reviewed that restaurant.

Example Join Operation

→ for all restaurants

```
db.restaurants.find().forEach(
```

```
    function(Object) {
```

```
        If (Object.grades.length > 1) {
```

```
            var user = db.Users.find(Object.name);
```

```
            if (user!=null) {
```

```
                printjson(Object.name);
```

```
                printjson(user);
```

```
            }
```

Manual ops

- decide how
to implement

```
        }
```

→ check if grade exists -

↑
look for
see who
reviewed

User,

User that has
restaurant in
"reviews"

)
- difficult, don't know
how data is stored

- deal w/ non-existent coll. / failures

Class Activity

- You have two collections in a MongoDB database. *marks* contains the list of students in a course with their marks and student number, and *students* contains the student number along with details such as first name, last name etc. How will you compute the join of these two collections (in JS code) from the Mongdb shell to list the student details along with the marks. You can assume the database is already loaded into the shell.

Solution to the activity

```
db.marks.find().forEach(  
    function( Object ) {  
        var st = db.students.find( {"student no":  
            Object.studentNo} );  
        if (st!=null) {  
            printjson(st);  
            printjson(Object.marks);  
        }  
        else {  
            print("No match found for " + Object.studentno);  
        }  
    }  
)
```

Outline

- Relational Databases (SQL-based)
- ACID semantics
- Non-traditional Databases (NoSQL)
- MongoDB Primer