

Class Activity – Load Balancing HTTP Requests



50

- Consider a cloud with 4 machines that accepts HTTP requests. Each requests triggers some “heavy” processing, and one server would not be enough to satisfy the demand – hence, we need to spread the requests between the 4 node servers.
- We will not be using a real could – you will be emulating this setup on your local machine.

Class Activity – Load Balancing HTTP Requests (2)



51

- `wasteCycles.js` is a Node program that serves the following HTTP requests:
 - `/cycles`: occupies the CPU for one second
 - (any string ending with `.html` or `.js`): returns the contents of the file (as you did in class)
 - Due to the single-threadness of Node, if you execute more than one request per second, the program will delay the processing of further requests
 - The program takes as input a port number to listen to.
- `waste.js` is a Node program that submits a request to `wasteCycles.js` every 800ms.
- Upon running `waste.js`, you will notice that the processing time goes well above 1000ms

Your task:



52

- You should launch four instances of `wasteCycles.js` to emulate four cloud servers
 - `node wasteCycles.js 8080`, `node wasteCycles.js 8081`, `node wasteCycles.js 8082`, `node wasteCycles.js 8083`
- Implement `loadBalancer.js` , a Node program accepts requests on port 8079 and submits them in a round-robin fashion to one of the node servers (`wasteCycles`)
- To test, modify the port number in `waste.js` to be 8079
- If your implementation is correct, the processing time should be stable around 1000ms